HMMs and the forward-backward algorithm

Ramesh Sridharan[*]

These notes give a short review of Hidden Markov Models (HMMs) and the forward-backward algorithm. They're written assuming familiarity with the sum-product belief propagation algorithm, but should be accessible to anyone who's seen the fundamentals of HMMs before.

The notation here is borrowed from Introduction to Probability by Bertsekas & Tsitsiklis: random variables are represented with capital letters, values they take are represented with lowercase letters, p_X represents a probability distribution for random variable X, and $p_X(x)$ represents the probability of value x (according to p_X).

Hidden Markov Models

Figure 1 shows the (undirected) graphical model for HMMs. Here's a quick recap of the important facts:

Figure 1: An undirected graphical model for the HMM. Connections between nodes indicate dependence.

- We observe Y_1 through Y_n, which we model as being observed from hidden states X_1 through X_n.

- Any particular state variable X_k depends only on X_{k-1} (what came before it), X_{k+1} (what comes after it), and Y_k (the observation associated with it).

- The goal of the forward-backward algorithm is to find the conditional distribution over hidden states given the data.

- In order to specify an HMM, we need three pieces:


[*]Contact: rameshvs@csail.mit.edu

Figure 2: A visualization of the forward and backward messages. Each message is a table that indicates what the node at the start point believes about the node at the end point.

- A *transition distribution*, $p_{X_{k+1}|X_k}(x_{k+1}|x_k) = W(x_{k+1}|x_k)$ [1], which describes the distribution for the next state given the current state. This is often represented as a matrix that we'll call A. Rows of A correspond to the current state, columns correspond to the next state, and each entry corresponds to the transition probability. So, the entry at row i and column j, A_{ij}, is $p_{X_{k+1}|X_k}(j|i)$, or equivalently $W(j|i)$.

- An *observation distribution* (also called an "emission distribution") $p_{Y_k|X_k}(y_k|x_k) = p_{Y|X}(y_k|x_k)$ [2], which describes the distribution for the output given the current state. We'll represent this with matrix B. Here, rows correspond to the current state, and columns correspond to the observation. So, $B_{ij} = p_{Y|X}(j|i)$: the probability of observing output j from state i is B_{ij}. Since the number of possible observations isn't necessarily the same as the number of possible states, B won't necessarily be square.

- An *initial state distribution* p_{X_1}, which describes the starting distribution over states. We'll represent this with a vector called π_0, where item i in the vector represents $p_{X_1}(i)$.

- The forward-backward algorithm computes forward and backward messages as follows:

$$m_{(k-1)\to k}(x_k) = \sum_{x_{k-1}} \overbrace{m_{(k-2)\to(k-1)}(x_{k-1})}^{\text{prev. message}} \overbrace{p_{Y|X}(y_{k-1}|x_{k-1})}^{\text{observation term}} \overbrace{W(x_{k-1}|x_k)}^{\text{transition term}}$$

$$m_{(k+1)\to k}(x_k) = \sum_{x_{k+1}} \underbrace{m_{(k+2)\to(k+1)}(x_{k+1})}_{\text{prev. message}} \underbrace{p_{Y|X}(y_{k+1}|x_{k+1})}_{\text{observation term}} \underbrace{W(x_k|x_{k+1})}_{\text{transition term}}$$

These messages are illustrated in Figure 2. The first forward message $m_{0\to 1}(x_1)$ is initialized to $\pi_0(x_1) = p_{X_1}(x_1)$. The first backward message $m_{(n+1)\to n}(x_n)$ is initialized to uniform (this is equivalent to not including it at all).

Figure 3 illustrates the computation of one forward message $m_{2\to 3}(x_3)$.

- To obtain a marginal distribution for a particular state given all the observations, $p_{X_k|Y_1,\dots,Y_n}$, we simply multiply the incoming messages together with the observation

---

[1] We're only going to worry about *homogeneous* Markov chains, where the transition distribution doesn't change over time: that's why our W and A notations only depend on the values and not the timepoints.

[2] Once again, we'll focus on Markov chains where the emission distribution is the same for every state.

2

term, and then normalize:

$$p_{X_k|Y_1,\ldots,Y_n}(x_k|y_1,\ldots,y_n) \propto m_{(k-1)\to k}(x_k)m_{(k+1)\to k}(x_k)p_{Y|X}(y_k|x_k)$$

Here, the symbol $\propto$ means "is proportional to", and indicates that we have to normalize at the end so that the answer sums to 1.

- Traditionally, the forward-backward algorithm computes a slightly different set of messages. The forward message α_k represents a message from $k-1$ to k that includes $p_{Y|X}(y_k|x_k)$, and the backward message β_k represents a message from $k+1$ to k identical to $m_{(k+1)\to k}$ above.

$$\alpha_k(x_k) = \overbrace{p_{Y|X}(y_k|x_k)}^{\text{observation term}} \sum_{x_{k-1}} \overbrace{\alpha_{k-1}(x_{k-1})}^{\text{prev. message}} \overbrace{W(x_{k-1}|x_k)}^{\text{transition term}}$$

$$\beta_k(x_k) = \sum_{x_{k+1}} \underbrace{\beta_{k+1}(x_{k+1})}_{\text{prev. message}} \underbrace{p_{Y|X}(y_{k+1}|x_{k+1})}_{\text{observation term}} \underbrace{W(x_k|x_{k+1})}_{\text{transition term}}$$

These messages have a particularly nice interpretation as probabilities:

$$\alpha_k(x_k) = p_{Y_1,Y_2,\ldots,Y_k,X_k}(y_1, y_2, \ldots, y_k, x_k)$$
$$\beta_k(x_k) = p_{Y_{k+1},Y_{k+2},\ldots,Y_n|X_k}(y_{k+1}, y_{k+2}, \ldots, y_n|x_k)$$

The initial forward α message is initialized to $\alpha_1(x_1) = p_{X_1}(x_1)p_{Y|X}(y_1|x_1)$. To obtain a marginal distribution, we simply multiply the messages together and normalize:

$$p_{X_k|Y_1,\ldots,Y_n}(x_k|y_1,\ldots,y_n) \propto \alpha_k(x_k)\beta_k(x_k)$$

Example

Suppose you send a robot to Mars. Unfortunately, it gets stuck in a canyon while landing and most of its sensors break. You know the canyon has 3 areas. Areas 1 and 3 are sunny and hot, while Area 2 is cold. You decide to plan a rescue mission for the robot from Area 3, knowing the following things about the robot:


Figure 3: An illustration of how to compute $m_{2\to3}(x_3)$. In order for node 2 to summarize its belief about X_3, it must incorporate the previous message $m_{1\to2}(x_2)$, its observation $p_{Y|X}(y_2|x_2)$, and the relationship $W(x_3|x_2)$ between X_2 and X_3.

- Every hour, it tries to move forward by one area (i.e. from Area 1 to Area 2, or Area 2 to Area 3). It succeeds with probability 0.75 and fails with probability 0.25. If it fails, it stays where it is. If it is in Area 3, it always stays there (and waits to be rescued).

- The temperature sensor still works. Every hour, we get a binary reading telling us whether the robot's current environment is hot or cold.

- We have no idea where the robot initially got stuck.

Solution:

(a) Construct an HMM for this problem: define a transition matrix A, an observation matrix B, and an initial state distribution π_0.

(b) Suppose we observe the sequence (hot, cold, hot). First, before doing any computation, determine the sequence of locations. Then, compute the forward and backward messages, and determine the distribution for the second state using the messages. Do your answers match up?

(a) We'll start with the transition matrix. Remember that each row corresponds to the current state, and each column corresponds to the next state. We'll use 3 states, each corresponding to an area.

- If the robot is in Area 1, it stays where it is with probability 0.25, moves to Area 2 with probability 0.75, and can't move to Area 3.

- Similarly, if the robot is in Area 2, it stays where it is with probability 0.25, can't move back to Area 1, and moves to Area 3 with probability 0.75.

- If the robot is in Area 3, it always stays in Area 3.

Each item above gives us one row of A. Putting it all together, we obtain

$$A = \begin{array}{c} 1 \\ 2 \\ 3 \end{array} \overset{\displaystyle \begin{array}{ccc} 1 & 2 & 3 \end{array}}{\begin{pmatrix} 0.25 & 0.75 & 0 \\ 0 & 0.25 & 0.75 \\ 0 & 0 & 1 \end{pmatrix}}$$

Next, let's look at the observation matrix. There are two possible observations, hot and cold. Areas 1 and 3 always produce "hot" readings while Area 2 always produces a "cold" reading:

$$B = \begin{array}{c} 1 \\ 2 \\ 3 \end{array} \overset{\displaystyle \begin{array}{cc} \text{hot} & \text{cold} \end{array}}{\begin{pmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 0 \end{pmatrix}}$$

4

Last but not least, since we have no idea where the robot starts, our initial state distribution will be uniform:

$$\pi_0 = \begin{matrix} 1 \\ 2 \\ 3 \end{matrix} \begin{pmatrix} 1/3 \\ 1/3 \\ 1/3 \end{pmatrix}$$

(b) Before doing any computation, we see that the sequence (hot,cold,hot) could only have been observed from the hidden state sequence (1,2,3). Make sure you convince yourself this is true before continuing!

We'll start with the forward messages.

$$m_{1 \to 2} = \sum_{x_1} \underbrace{m_{0 \to 1}(x_1) p_{Y|X}(y_1|x_1)}_{\text{depends only on } x_1 \text{ and } y_1} \psi(x_1, x_2)$$

The output message should have three different possibilities, one for each value of x_2. We can therefore represent it as a vector indexed by x_2:

$$\begin{pmatrix} \cdot \\ \cdot \\ \cdot \end{pmatrix} \begin{matrix} \text{value for } x_2 = 1 \\ \text{value for } x_2 = 2 \\ \text{value for } x_2 = 3 \end{matrix}$$

For each term in the sum (i.e., each possible value of x_1):

- $m_{0 \to 1}$ comes from from the initial distribution. Normally it would come from the previous message, but our first forward message is always set to initial state distribution.

- $p_{Y|X}(y_1|x_1)$ comes from the column of B corresponding to our observation $y_1 = $ hot.

- ψ comes from a row of A: we are fixing x_1 and asking about possible values for x_2, which corresponds exactly to the transition distributions given in the rows of A (remember that the rows of A correspond to the current state and the columns correspond to the next state).

So, we obtain

$$m_{1 \to 2} = \overbrace{\frac{1}{3} \cdot 1 \cdot \begin{pmatrix} .25 \\ .75 \\ 0 \end{pmatrix}}^{x_1=1} + \overbrace{\frac{1}{3} \cdot 0 \cdot \begin{pmatrix} 0 \\ .25 \\ .75 \end{pmatrix}}^{x_1=2} + \overbrace{\frac{1}{3} \cdot 1 \cdot \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}}^{x_1=3}$$

$$\propto \begin{pmatrix} 1 \\ 3 \\ 4 \end{pmatrix}$$

Since our probabilities are eventually computed by multiplying messages and normalizing, we can arbitrary renormalize at any step to make the computation easier.

For the second message, we perform a similar computation:

$$m_{2\to3} = \sum_{x_2} m_{1\to2}(x_2)\tilde{\phi}(x_2)\psi(x_2, x_3)$$

$$= 1 \cdot 0 \cdot \overbrace{\begin{pmatrix} .25 \\ .75 \\ 0 \end{pmatrix}}^{x_2=1} + 3 \cdot 1 \cdot \overbrace{\begin{pmatrix} 0 \\ .25 \\ .75 \end{pmatrix}}^{x_2=2} + 4 \cdot 0 \cdot \overbrace{\begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}}^{x_2=3}$$

$$\propto \begin{pmatrix} 0 \\ 1 \\ 3 \end{pmatrix}$$

The backwards messages are computed using a similar formula:

$$m_{3\to2} = \sum_{x_3} \underbrace{m_{4\to3}(x_3)\tilde{\phi}(x_3)}_{\text{depends only on } x_3} \psi(x_2, x_3)$$

The first backwards message, $m_{4\to3}(x_3)$, is always initialized to uniform since we have no information about what the last state should be. Note that this is equivalent to not including that term at all.

For each value of x_3, the transition term $\psi(x_2, x_3)$ is now drawn from a *column* of A, since we are interested in the probability of arriving at x_3 from each possible state for x_2. We compute the messages as:

$$m_{3\to2} = 1 \cdot \overbrace{\begin{pmatrix} .25 \\ 0 \\ 0 \end{pmatrix}}^{x_3=1} + 0 \cdot \overbrace{\begin{pmatrix} .75 \\ .25 \\ 0 \end{pmatrix}}^{x_3=2} + 1 \cdot \overbrace{\begin{pmatrix} 0 \\ .75 \\ 1 \end{pmatrix}}^{x_3=3}$$

$$\propto \begin{pmatrix} 1 \\ 3 \\ 4 \end{pmatrix}$$

Similarly, the second backwards message is:

$$m_{2\to1} = 1 \cdot 0 \cdot \overbrace{\begin{pmatrix} .25 \\ 0 \\ 0 \end{pmatrix}}^{x_2=1} + 3 \cdot 1 \cdot \overbrace{\begin{pmatrix} .75 \\ .25 \\ 0 \end{pmatrix}}^{x_2=2} + 4 \cdot 0 \cdot \overbrace{\begin{pmatrix} 0 \\ .75 \\ 1 \end{pmatrix}}^{x_2=3}$$

$$\propto \begin{pmatrix} 3 \\ 1 \\ 0 \end{pmatrix}$$

Notice from the symmetry of the problem that our forwards messages and backwards messages were the same.

To compute the marginal distribution for X_2 given the data, we multiply the messages and the observation:

$$p_{X_2|Y_1,\ldots,Y_n}(x_2|y_1,\ldots,y_n) \propto m_{1\to 2}(x_2)m_{3\to 2}(x_2)\tilde{\phi}(x_2)$$

$$\propto \begin{pmatrix} 1 \\ 3 \\ 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 3 \\ 4 \end{pmatrix} \cdot \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}$$

$$= \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}$$

Notice that in this case, because of our simplified observation model, the observation "cold" allowed us to determine the state. This matches up with our earlier conclusion that the robot must have been in Area 2 during the second hour.

If we were to compute α messages, we would start with our initial message, α_1:

$$\alpha_1(x_1) = p_{X_1}(x_1)p_{Y|X}(y_1|x_1) = \begin{pmatrix} 1/3 \\ 0 \\ 1/3 \end{pmatrix}$$

The first real message is computed as follows:

$$\alpha_2 = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \cdot \left(\overbrace{\cdot 1/3 \cdot \begin{pmatrix} .25 \\ .75 \\ 0 \end{pmatrix}}^{x_1=1} + \overbrace{\cdot 0 \cdot \begin{pmatrix} 0 \\ .25 \\ .75 \end{pmatrix}}^{x_1=2} + \overbrace{\cdot 1/3 \cdot \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}}^{x_1=3} \right)$$

$$\propto \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}$$

The second message is similar:

$$\alpha_3 = \begin{pmatrix} 1 \\ 0 \\ 1 \end{pmatrix} \cdot \left(\overbrace{\cdot 0 \cdot \begin{pmatrix} .25 \\ .75 \\ 0 \end{pmatrix}}^{x_1=1} + \overbrace{\cdot 1 \cdot \begin{pmatrix} 0 \\ .25 \\ .75 \end{pmatrix}}^{x_1=2} + \overbrace{\cdot 0 \cdot \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}}^{x_1=3} \right)$$

$$\propto \begin{pmatrix} 1 \\ 0 \\ 1 \end{pmatrix}$$

The β messages would be identical to our backwards messages computed earlier.